

```
# TEMEL
import numpy as np
import pandas as pd

# GÖRSELLEŞTİRME
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# UYARILARI KAPAT
import warnings
warnings.filterwarnings("ignore")

# MAKİNE ÖĞRENMESİ - TEMEL
from sklearn.model_selection import train_test_split, GridSearchCV,
learning_curve
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.pipeline import Pipeline

# MODELLER
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier

# DEĞERLENDİRME METRİKLERİ
from sklearn.metrics import (
    accuracy_score,
    confusion_matrix,
    ConfusionMatrixDisplay,
    classification_report,
    roc_auc_score,
    RocCurveDisplay
)
# OPSİYONEL GÜÇLÜ MODELLER
# !pip install catboost
# from catboost import CatBoostClassifier

# !pip install xgboost
# from xgboost import XGBClassifier

df =
pd.read_csv('~/Users/tugcekizilkoca/Desktop/veri-madenciligi/datasets/
heart.csv')

# Returns number of rows and columns of the dataset
df.shape
(303, 14)
```

```

# Returns an object with all of the column headers
df.columns

Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg',
'thalach',
       'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
      dtype='object')

# Returns different datatypes for each columns (float, int, string,
bool, etc.)
df.dtypes

age          int64
sex          int64
cp           int64
trestbps    int64
chol          int64
fbs           int64
restecg     int64
thalach      int64
exang          int64
oldpeak      float64
slope          int64
ca            int64
thal           int64
target         int64
dtype: object

# Returns the first x number of rows when head(x). Without a number it
returns 5
df.head()

   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak
slope \
0    63    1    3        145   233    1        0      150      0     2.3
0
1    37    1    2        130   250    0        1      187      0     3.5
0
2    41    0    1        130   204    0        0      172      0     1.4
2
3    56    1    1        120   236    0        1      178      0     0.8
2
4    57    0    0        120   354    0        1      163      1     0.6
2

   ca  thal  target
0    0    1      1
1    0    2      1
2    0    2      1

```

```
3   0    2    1  
4   0    2    1
```

```
# Returns the last x number of rows when tail(x). Without a number it  
returns 5  
df.tail()
```

```
      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  
oldpeak \\  
298    57    0    0        140    241    0        1       123      1  
0.2  
299    45    1    3        110    264    0        1       132      0  
1.2  
300    68    1    0        144    193    1        1       141      0  
3.4  
301    57    1    0        130    131    0        1       115      1  
1.2  
302    57    0    1        130    236    0        0       174      0  
0.0
```

```
      slope  ca  thal  target  
298      1    0    3      0  
299      1    0    3      0  
300      1    2    3      0  
301      1    1    3      0  
302      1    1    2      0
```

```
# Returns true for a column having null values, else false  
df.isnull().any()
```


```
age        False  
sex        False  
cp         False  
trestbps  False  
chol        False  
fbs         False  
restecg    False  
thalach    False  
exang      False  
oldpeak    False  
slope      False  
ca          False  
thal        False  
target     False  
dtype: bool
```


```
# Returns basic statistics on numeric columns  
df.describe().T
```


	count	mean	std	min	25%	50%	75%
max							


age	303.0	54.366337	9.082101	29.0	47.5	55.0	61.0
77.0							
sex	303.0	0.683168	0.466011	0.0	0.0	1.0	1.0
1.0							
cp	303.0	0.966997	1.032052	0.0	0.0	1.0	2.0
3.0							
trestbps	303.0	131.623762	17.538143	94.0	120.0	130.0	140.0
200.0							
chol	303.0	246.264026	51.830751	126.0	211.0	240.0	274.5
564.0							
fbs	303.0	0.148515	0.356198	0.0	0.0	0.0	0.0
1.0							
restecg	303.0	0.528053	0.525860	0.0	0.0	1.0	1.0
2.0							
thalach	303.0	149.646865	22.905161	71.0	133.5	153.0	166.0
202.0							
exang	303.0	0.326733	0.469794	0.0	0.0	0.0	1.0
1.0							
oldpeak	303.0	1.039604	1.161075	0.0	0.0	0.8	1.6
6.2							
slope	303.0	1.399340	0.616226	0.0	1.0	1.0	2.0
2.0							
ca	303.0	0.729373	1.022606	0.0	0.0	0.0	1.0
4.0							
thal	303.0	2.313531	0.612277	0.0	2.0	2.0	3.0
3.0							
target	303.0	0.544554	0.498835	0.0	0.0	1.0	1.0
1.0							


```
df.hist(figsize=(15,15))
plt.tight_layout()
plt.show()
```












```

sns.countplot(x='target', data=df)
plt.xlabel("Target")
plt.ylabel("Count")
plt.show()


```



```
# Selecting correlated features using Heatmap

# Get correlation of all the features of the dataset
corr_matrix = df.corr()
top_corr_features = corr_matrix.index

# Plotting the heatmap
plt.figure(figsize=(20,20))
sns.heatmap(df[top_corr_features].corr(), annot=True, cmap='RdYlGn')
plt.title("Correlation Heatmap")
plt.show()
```



```
dataset = pd.get_dummies(df, columns=['sex', 'cp', 'fbs', 'restecg',
                                         'exang', 'slope', 'ca', 'thal'])

dataset.columns

Index(['age', 'trestbps', 'chol', 'thalach', 'oldpeak', 'target',
       'sex_0',
       'sex_1', 'cp_0', 'cp_1', 'cp_2', 'cp_3', 'fbs_0', 'fbs_1',
       'restecg_0',
```

```

        'restecg_1', 'restecg_2', 'exang_0', 'exang_1', 'slope_0',
'slope_1',
        'slope_2', 'ca_0', 'ca_1', 'ca_2', 'ca_3', 'ca_4', 'thal_0',
'thal_1',
        'thal_2', 'thal_3'],
      dtype='object')

from sklearn.preprocessing import StandardScaler
standScaler = StandardScaler()
columns_to_scale = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
dataset[columns_to_scale] =
standScaler.fit_transform(dataset[columns_to_scale])

dataset.head()

      age  trestbps      chol      thalach      oldpeak  target  sex_0
sex_1 \
0  0.952197  0.763956 -0.256334  0.015443  1.087338      1  False
True
1 -1.915313 -0.092738  0.072199  1.633471  2.122573      1  False
True
2 -1.474158 -0.092738 -0.816773  0.977514  0.310912      1   True
False
3  0.180175 -0.663867 -0.198357  1.239897 -0.206705      1  False
True
4  0.290464 -0.663867  2.082050  0.583939 -0.379244      1   True
False

      cp_0  cp_1  ...  slope_2  ca_0  ca_1  ca_2  ca_3  ca_4
thal_0 \
0  False  False  ...  False  True  False  False  False  False
False
1  False  False  ...  False  True  False  False  False  False
False
2  False  True  ...  True  True  False  False  False  False
False
3  False  True  ...  True  True  False  False  False  False
False
4  True  False  ...  True  True  False  False  False  False
False

      thal_1  thal_2  thal_3
0    True  False  False
1  False  True  False
2  False  True  False
3  False  True  False
4  False  True  False

[5 rows x 31 columns]

```

```

# Splitting the dataset into dependent and independent features
X = dataset.drop('target', axis=1)
y = dataset['target']

#KNN

# Importing essential libraries
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score

# Finding the best accuracy for knn algorithm using cross_val_score
knn_scores = []
for i in range(1, 21):
    knn_classifier = KNeighborsClassifier(n_neighbors=i)
    cvs_scores = cross_val_score(knn_classifier, X, y, cv=10)
    knn_scores.append(round(cvs_scores.mean(),3))

plt.figure(figsize=(20,15))


# skor grafiği
plt.plot(range(1,21), knn_scores, color='red')

# her noktanın üzerine skor yaz
for i, score in enumerate(knn_scores, start=1):
    plt.text(i, score, f"{score:.2f}")

# eksen ayarları
plt.xticks(range(1,21))
plt.xlabel('Number of Neighbors (K)')
plt.ylabel('Scores')
plt.title('KNN Scores for Different K Values')

plt.show()

```



```

# Training the knn classifier model with k value as 12
knn_classifier = KNeighborsClassifier(n_neighbors=12)
cvs_scores = cross_val_score(knn_classifier, X, y, cv=10)
print("KNeighbours Classifier Accuracy with K=12 is: {}"
%.format(round(cvs_scores.mean(), 4)*100))

KNeighbours Classifier Accuracy with K=12 is: 84.48%

# Importing essential libraries
from sklearn.tree import DecisionTreeClassifier

# Finding the best accuracy for decision tree algorithm using
cross_val_score
decision_scores = []
for i in range(1, 11):
    decision_classifier = DecisionTreeClassifier(max_depth=i)
    cvs_scores = cross_val_score(decision_classifier, X, y, cv=10)
    decision_scores.append(round(cvs_scores.mean(),3))


# Plotting the results of decision_scores
plt.figure(figsize=(20,15))
plt.plot([i for i in range(1, 11)], decision_scores, color = 'red')

```

```

for i in range(1,11):
    plt.text(i, decision_scores[i-1], (i, decision_scores[i-1]))
plt.xticks([i for i in range(1, 11)])
plt.xlabel('Depth of Decision Tree (N)')
plt.ylabel('Scores')
plt.title('Decision Tree Classifier scores for different depth values')
plt.show()

```



```

# Training the decision tree classifier model with max_depth value as 3
decision_classifier = DecisionTreeClassifier(max_depth=3)
cvs_scores = cross_val_score(decision_classifier, X, y, cv=10)
print("Decision Tree Classifier Accuracy with max_depth=3 is: {}%".format(round(cvs_scores.mean(), 4)*100))

Decision Tree Classifier Accuracy with max_depth=3 is: 78.51%

# Importing essential libraries
from sklearn.ensemble import RandomForestClassifier

```

```

# Finding the best accuracy for random forest algorithm using
cross_val_score
forest_scores = []
for i in range(10, 101, 10):
    forest_classifier = RandomForestClassifier(n_estimators=i)
    cvs_scores = cross_val_score(forest_classifier, X, y, cv=5)
    forest_scores.append(round(cvs_scores.mean(),3))

# Plotting the results of forest_scores
plt.figure(figsize=(20,15))
plt.plot([n for n in range(10, 101, 10)], forest_scores, color =
'red')
for i in range(1,11):
    plt.text(i*10, forest_scores[i-1], (i*10, forest_scores[i-1]))
plt.xticks([i for i in range(10, 101, 10)])
plt.xlabel('Number of Estimators (N)')
plt.ylabel('Scores')
plt.title('Random Forest Classifier scores for different N values')

Text(0.5, 1.0, 'Random Forest Classifier scores for different N
values')

# Training the random forest classifier model with n value as 90
forest_classifier = RandomForestClassifier(n_estimators=90)
cvs_scores = cross_val_score(forest_classifier, X, y, cv=5)
print("Random Forest Classifier Accuracy with n_estimators=90 is: {}%
".format(round(cvs_scores.mean(), 4)*100))

Random Forest Classifier Accuracy with n_estimators=90 is: 81.83%

# Hedef değişken ismini burada ayarla:
TARGET_COL = 'target' # kalp hastalığı datası için
# TARGET_COL = 'Completed' # kurs tamamlama datası için

# X, y ayrımı
X = df.drop(columns=[TARGET_COL])
y = df[TARGET_COL]

# Train-test ayırma
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

X_train.shape, X_test.shape
((242, 13), (61, 13))

```